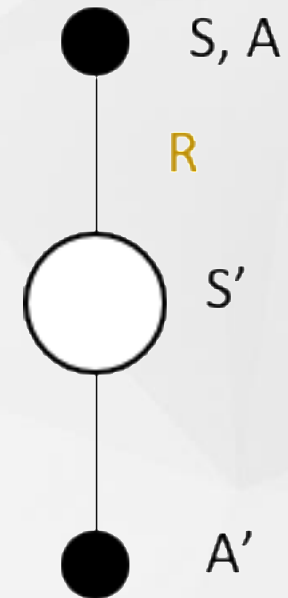


## Chapter 06. 스스로 전략을 짜는 강화학습 ( Reinforcement Learning )

# Sarsa , Q- learning

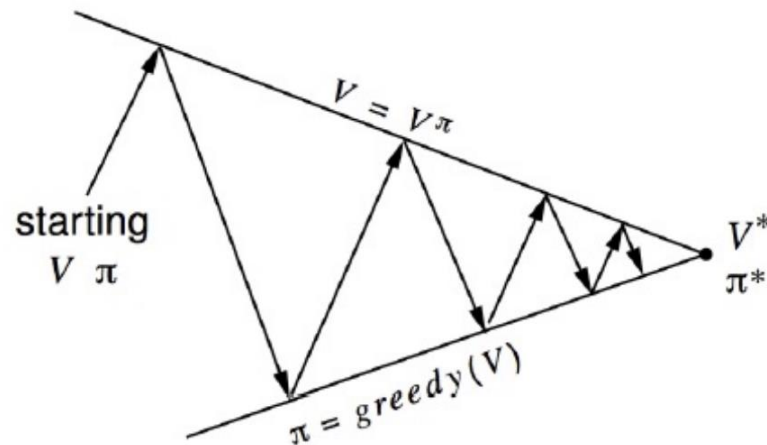


# Monte Carlo Method

## Monte-Carlo Control

**Policy Iteration** = ( policy evaluation + policy improvement )

**Monte-Carlo policy Iteration** = ( MC policy evaluation + policy improvement )



Policy evaluation Monte-Carlo policy evaluation,  $V = v_\pi$ ?

Policy improvement Greedy policy improvement?

**Problem 1 . Value Function ->**  
**MDD**

$$v(s) = \mathcal{R}_s + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'} v(s')$$

**Problem 2 . Greedy policy improvement**

**-> Local Optimum**

# Monte Carlo Method

## Problem 1 . Value Function -> MDP

- Greedy policy improvement over  $V(s)$  requires model of MDP

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} \mathcal{R}_s^a + \mathcal{P}_{ss'}^a V(s')$$

- Greedy policy improvement over  $Q(s, a)$  is model-free

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a)$$

## Problem 2 . Greedy policy improvement -> Local Optimum

$$\pi(s) \doteq \operatorname{argmax}_a q(s, a).$$

$$\begin{aligned} q_{\pi_k}(s, \pi_{k+1}(s)) &= q_{\pi_k}(s, \operatorname{argmax}_a q_{\pi_k}(s, a)) \\ &= \max_a q_{\pi_k}(s, a) \\ &\geq q_{\pi_k}(s, \pi_k(s)) \\ &\geq v_{\pi_k}(s). \end{aligned}$$

- Simplest idea for ensuring continual exploration
- All  $m$  actions are tried with non-zero probability
- With probability  $1 - \epsilon$  choose the greedy action
- With probability  $\epsilon$  choose an action at random

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon & \text{if } a^* = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a) \\ \epsilon/m & \text{otherwise} \end{cases}$$

# Temporal Difference

에피소드마다가 아니라 매 타임스텝마다 가치함수를 업데이트

$$\text{MC : } V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

$$\text{TD(0) : } V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

Input : the policy  $\pi$  to be evaluated

Initialize  $V(s)$  arbitrarily (e.g.,  $V(s) = 0, \forall s \in S^+$ )

Repeat (for each episode):

    Initialize  $S$

    Repeat (for each step of episode):

$A \leftarrow$  action given by  $\pi$  for  $S$

        Take action  $A$ ; observe reward,  $R$ , and next state,  $S'$

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

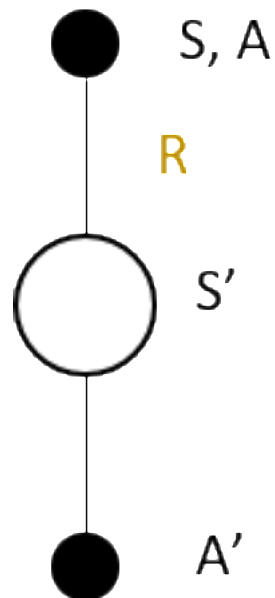
$S \leftarrow S'$

    until  $S$  is terminal

# Temporal Difference Control

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma Q(S', A') - Q(S, A))$$

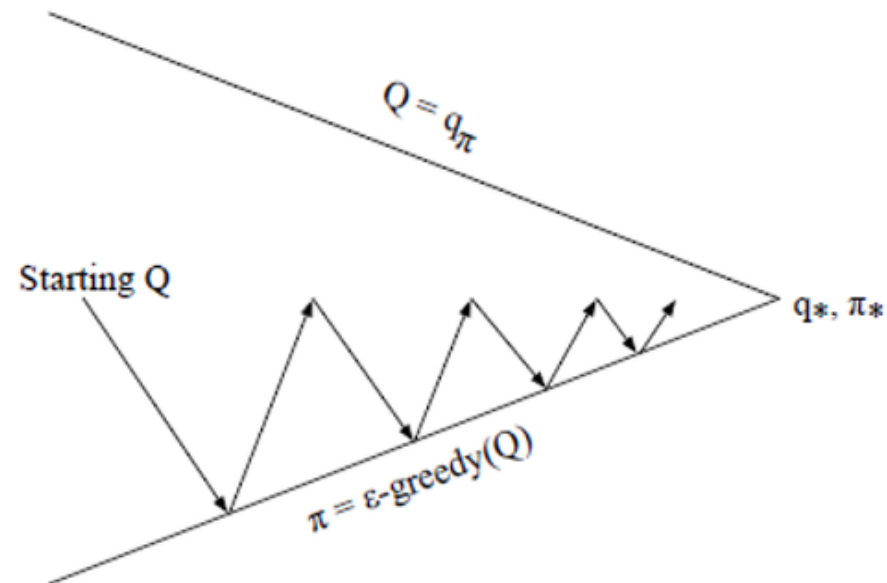


$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

$[S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}]$ 을 하나의 샘플로 사용하기 때문에 **SARSA**라고 합니다.  
앞으로는 시간차 제어가 아닌 살사라고 부르겠습니다.

# SARSA

## On-Policy Control With Sarsa



Every **time-step**:

Policy evaluation **Sarsa**,  $Q \approx q_\pi$

Policy improvement  $\epsilon$ -greedy policy improvement

# SARSA

Initialize  $Q(s, a)$ ,  $\forall s \in S, a \in A(s)$ , arbitrarily and  $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

    Repeat (for each step of episode):

        Take action  $A$ ; observe  $R, S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

    until  $S$  is terminal

# N- step SARSA

## $n$ -Step Sarsa

- Consider the following  $n$ -step returns for  $n = 1, 2, \infty$ :

$$\begin{array}{ll} n = 1 & \text{(Sarsa)} \quad q_t^{(1)} = R_{t+1} + \gamma Q(S_{t+1}) \\ n = 2 & q_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 Q(S_{t+2}) \\ & \vdots \\ n = \infty & \text{(MC)} \quad q_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T \end{array}$$

- Define the  $n$ -step Q-return

$$q_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q(S_{t+n})$$

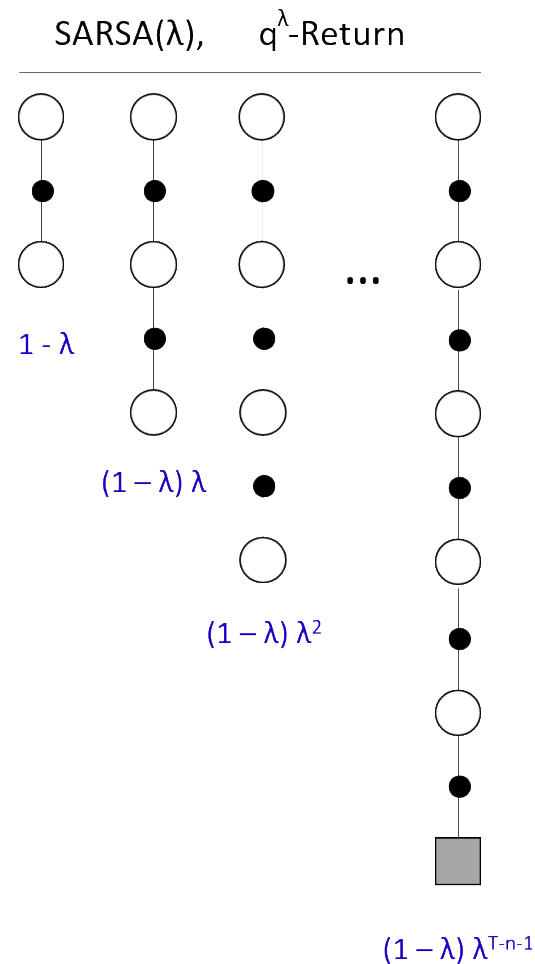
- $n$ -step Sarsa updates  $Q(s, a)$  towards the  $n$ -step Q-return

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left( q_t^{(n)} - Q(S_t, A_t) \right)$$



# SARSA( $\lambda$ )

## Forward Sarsa



$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(q_t^\lambda - Q(S_t, A_t))$$

$$\text{when } q^\lambda\text{-return, } q_t^\lambda = (1-\lambda) \sum_{n=1}^{\infty} \lambda^{n-1} q_t^{(n)}$$

# SARSA( $\lambda$ )

$$E_t(s, a) = \begin{cases} \gamma\lambda E_{t-1}(s, a) + 1 & \text{if } s = s_t, a = a_t \\ \gamma\lambda E_{t-1}(s, a) & \text{otherwise} \end{cases}$$

## Backward Sarsa

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \delta_t E_t(s, a)$$

when  $\delta_t = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)$  (TD error)

Initialize  $Q(s, a)$  arbitrarily,  $\forall s \in S, a \in A(s)$   
Repeat (for each episode):  
  Initialize  $S, A$   
  Repeat (for each step of episode):  
    Take action  $A$ ; observe  $R, S'$   
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
     $\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$   
     $E(S, A) \leftarrow E(S, A) + 1$   
    For all  $s \in S, a \in A(s)$ :  
       $Q(s, a) \leftarrow Q(s, a) + \alpha \delta E(s, a)$   
       $E(s, a) \leftarrow \gamma\lambda E(s, a)$   
     $S \leftarrow S'; A \leftarrow A'$   
  until  $S$  is terminal

# On / Off Policy

## On-policy :

학습하는 policy와 행동하는 policy가 반드시 같아야만 학습이 가능한 강화학습 알고리즘.

**ex) Sarsa**

on-policy의 경우 1번이라도 학습을 해서 policy improvement를 시킨 순간, 그 policy가 했던 과거의 experience들은 모두 사용이 불가능하다. 즉 매우 데이터 효율성이 떨어진다. 바로바로 exploration해서 학습하고 재사용이 불가능하다.

## Off-policy :

학습하는 policy와 행동하는 policy가 반드시 같지 않아도 학습이 가능한 알고리즘.

**ex) Q-learning**

off-policy는 현재 학습하는 policy가 과거에 했던 experience도 학습에 사용이 가능하고, 심지어는 해당 policy가 아니라 예를 들어 사람이 한 데이터로부터도 학습을 시킬 수가 있다.

# Off Policy

- Learn from observing humans or other agents
- Re-use experience generated from old policies  $\pi_1, \pi_2, \dots, \pi_{t-1}$
- Learn about ***Optimal Policy*** while following exploratory policy
- Learn about ***multiple policies*** while following one policy

# Q- Learning

<https://www.slideshare.net/DongMinLee32/part-2-91522217>

행동하는 정책

$\epsilon$ -탐욕 정책

학습하는 정책

다음 상태에서 어떤 행동을 할 때  
다음 상태의 **최대** 큐함수를  
현재 상태의 큐함수로 업데이트

# Q- Learning

<https://www.slideshare.net/DongMinLee32/part-2-91522217>

살사의 큐함수 업데이트

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$



큐러닝의 큐함수 업데이트

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a') - Q(S_t, A_t))$$

다음 상태에서 다음 행동을 해보는 것이 아니라  
다음 상태에서 **가장 큰 큐함수를 가지고 업데이트**

# Q- Learning

<https://www.slideshare.net/DongMinLee32/part-2-91522217>

살사의 필요한 샘플

$[S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}]$



큐러닝의 필요한 샘플

$[S_t, A_t, R_{t+1}, S_{t+1}]$

다음 상태에서 **가장 큰 큐함수만** 필요하기 때문에  
샘플도  $[S_t, A_t, R_{t+1}, S_{t+1}]$ 까지만 필요합니다.

# Q- Learning

- 1) 현재 state S 에서 behavior policy,  $\mu$ (e.g.  $\epsilon$ -greedy)에 따라 action A을 선택.
- 2) q-func.을 이용하여 다음 state S'에서의 action A'는  $\pi$ (e.g. greedy)에 따라 선택.

$$\pi(S_{t+1}) = \underset{a'}{\operatorname{argmax}} Q(S_{t+1}, a')$$

- 3) Q-learning의 target은 아래 식으로 도출.

$$\begin{aligned} R_{t+1} + \gamma Q(S_{t+1}, A') &= R_{t+1} + \gamma Q(S_{t+1}, \underset{a'}{\operatorname{argmax}} Q(S_{t+1}, a')) \\ &= R_{t+1} + \max_{a'} \gamma Q(S_{t+1}, a') \end{aligned}$$

- 4) 아래 식에 따라 q-func.을 update.

$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_{a'} Q(S', A') - Q(S, A))$$



# Q- Learning

<https://www.slideshare.net/CurtPark1/dqn-reinforcement-learning-from-basics-to-dqn>

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)].$$

Q-learning (off-policy TD control) for estimating  $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize  $S$

Loop for each step of episode:

Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

Take action  $A$ , observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

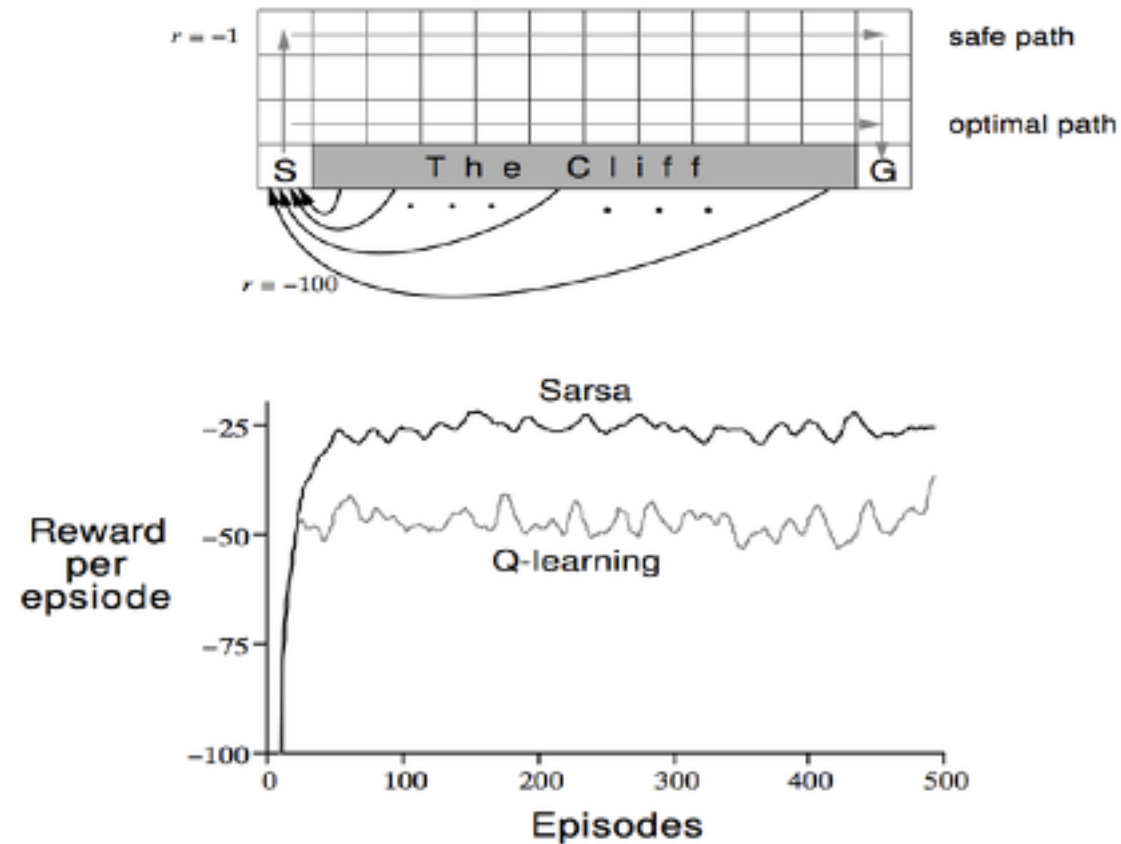
until  $S$  is terminal

Behavior policy로 동작

Target policy로 동작

# Sarsa & Q- Learning

## Cliff Walking Example



• *Thank you*